



## MODULE-2

### CRUD OPERATIONS

## → Module 1

- » Introduction and Overview
- » No SQL

## → Module 2

- » **CRUD Operations**
- » **CRUD Concerns**

## → Module 3

- » Schema Design and Data Modeling
- » Comparison with Relational Systems

## → Module 4

- » Administration
- » Backup and Recovery

## → Module 5

- » Scalability and Availability
- » Replication and Sharding

## → Module 6

- » Indexing and Aggregation Frame work
- » Performance Tuning

## → Module 7

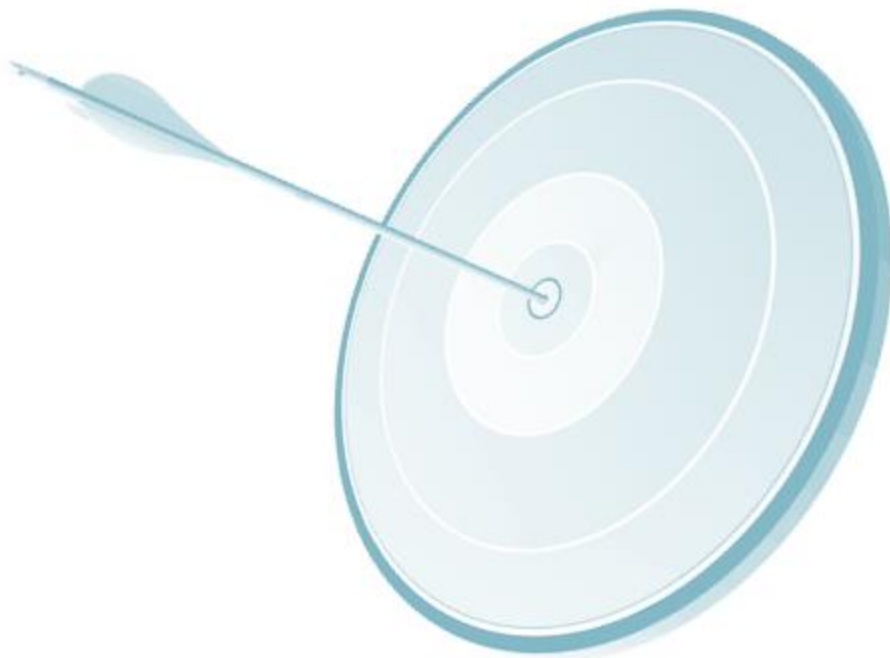
- » Application Engineering and MongoDB® Tools
- » Interface with Other Language

## → Module 8

- » Project, Additional Concepts and Cases Studies

At the end of this module, you will be able to

- Understand MongoDB®'s development and production architecture
- Understand read and write concepts of MongoDB®
- Understand how Journaling works
- Use mongo shell for CRUD operations
- Understand different MongoDB® data types



What are typical usage for MongoDB® ?



MongoDB® has a general-purpose design, making it appropriate for a large number of use cases. Examples include content management systems, mobile applications, gaming, e-commerce, analytics, archiving, and logging.



Where shouldn't MongoDB® used?



MongoDB® shouldn't be used for systems that require SQL, joins, and multi-object transactions.



Does MongoDB® handle caching?





Yes. MongoDB® keeps all of the most recently used data in RAM. If you have created indexes for your queries and your working data set fits in RAM, MongoDB® serves all queries from memory.



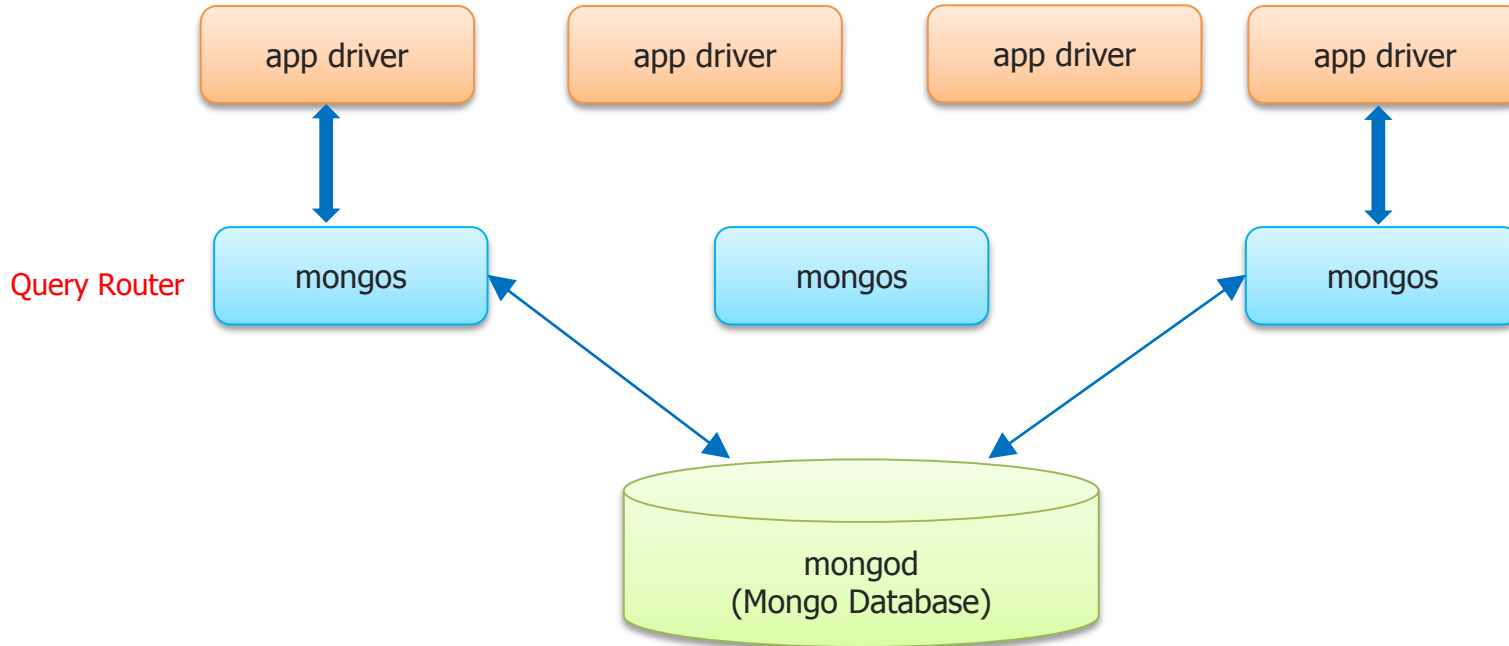
What is the difference between mongo and mongos?



**Mongo** - This javascript shell acts as a client and connects to a mongod or mongos instance for general database operations.

**Mongos** - This is a daemon which acts as the query router between the client (mongo or application layer) and the Sharded Cluster (one or multiple mongod).





# edureka!



# MongoDB® CRUD Introduction

→ MongoDB® provides rich functionalities for reading and manipulating data.

→ CRUD stands for Create, Read, Update and Delete.

C → Create

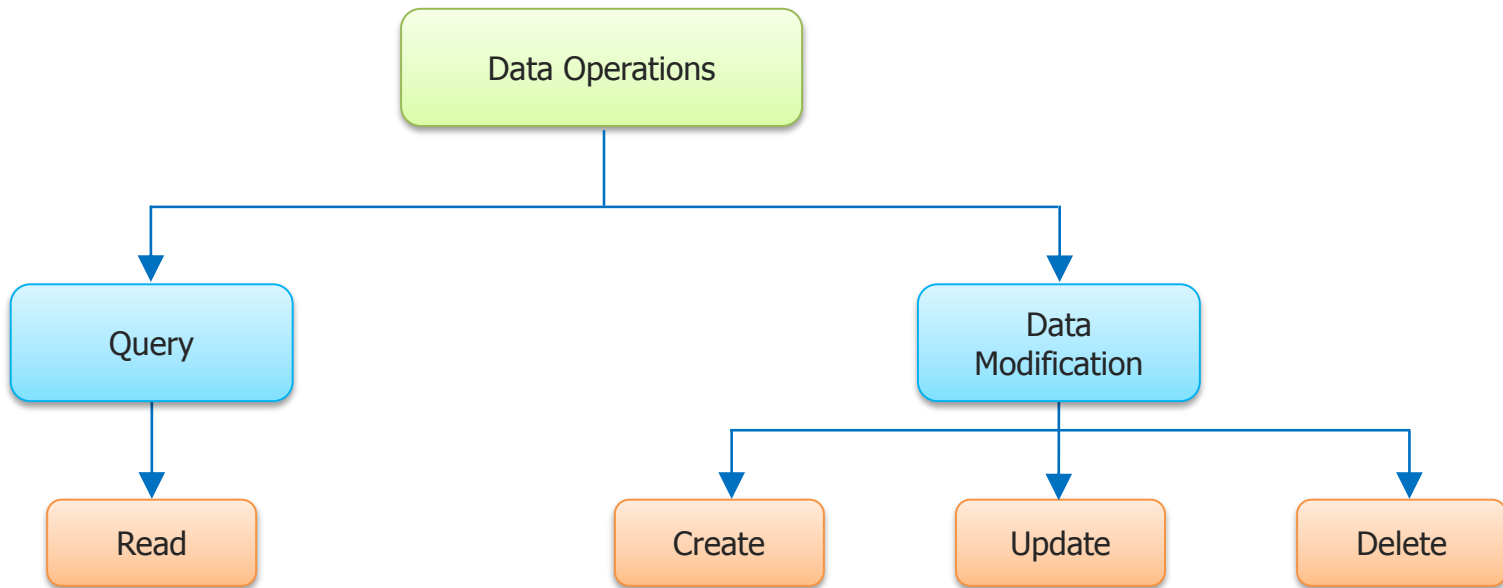
R → Read

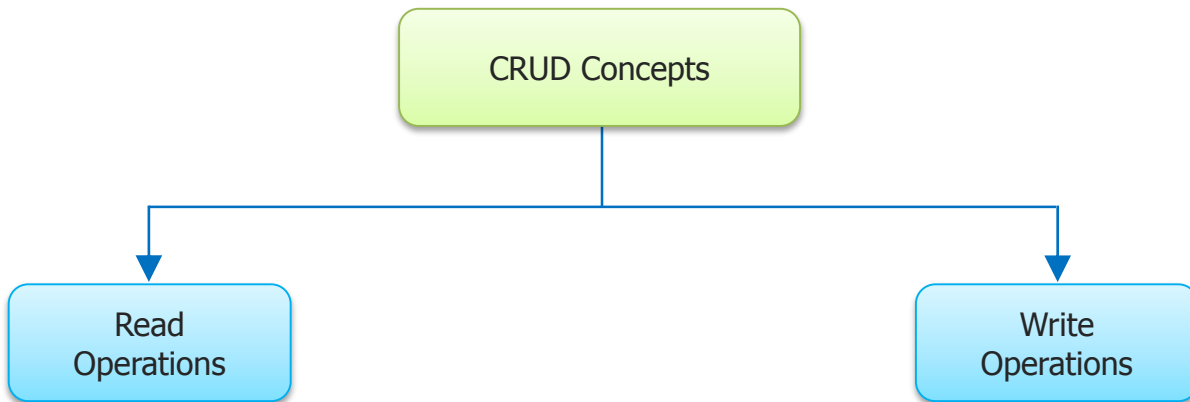
U → Update

D → Delete

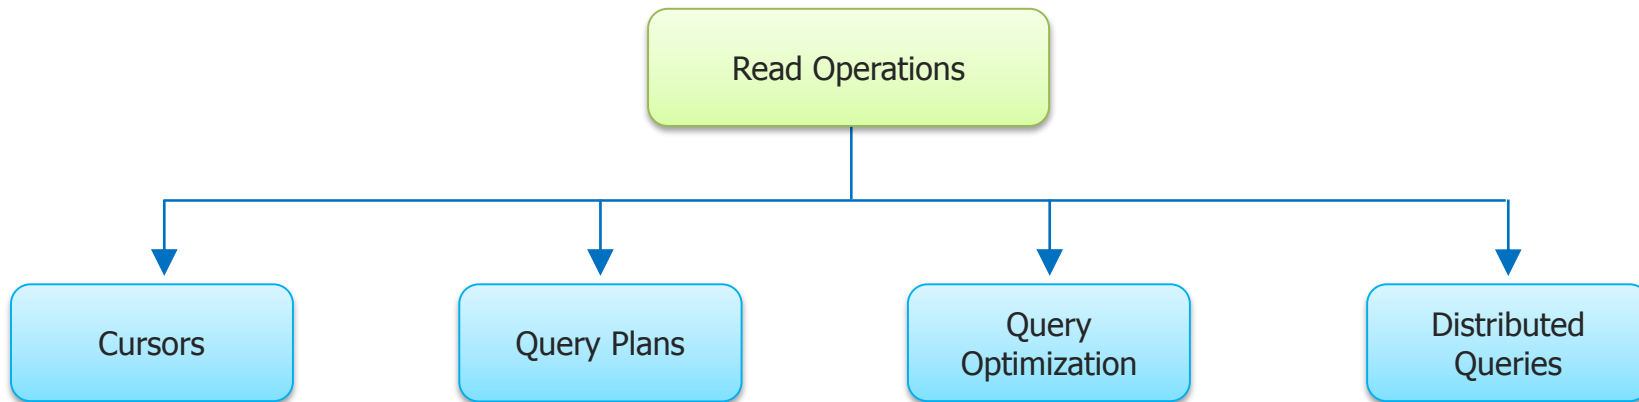


These terms are the basic stuffs for all interactions with the databases.









→ A MongoDB® Query :

db . <b>users</b> . find(	←	collection
{ age : { \$gt : 18 } },	←	query criteria
{ name : 1, address : 1 }	←	projection
) . limit (5)	←	cursor modifier

→ An Equivalent SQL query:

SELECT	_id , name , address	←	projection
FROM	users	←	table
WHERE	age > 18	←	select criteria
LIMIT	5	←	cursor modifier

- All queries in MongoDB® address a single collection.
- You can modify the query to impose [limits](#), [skips](#), and [sort orders](#).
- The order of documents returned by a query is not defined unless you specify a [sort\(\)](#).
- Operations that [modify existing documents](#) (i.e. updates) use the same query syntax as queries to select documents to update.
- In [aggregation](#) pipeline, the [\\$match](#) pipeline stage provides access to MongoDB® queries.
- MongoDB® provides a [db.collection.findOne\(\)](#) method as a special case of [find\(\)](#) that returns a single document.

→ Exclude One Field From a Result Set

```
db.records.find( { "user_id": { $lt: 42 } }, { history: 0 } )
```

→ Return Two Fields and the \_id Field

```
db.records.find( { "user_id": { $lt: 42 } }, { "name": 1, "email": 1 } )
```

→ Exclude \_id field

```
db.records.find( { "user_id": { $lt: 42 } }, { "_id": 0, "name": 1, "email": 1 } )
```

- Queries return iterable objects, called cursors, that hold the full result set of the query request.
- If the returned cursor is not assigned to a variable using the var keyword, then the cursor is automatically iterated up to 20 times.
- `db.collection.find()` method queries a collection and returns a cursor to the returning documents. To access the documents, you need to iterate the cursor.
- The `db.serverStatus()` will return cursor information. How many cursors are open, timed out etc. can be identified.

# Iterate through the Cursor

```
var myCursor = db.inventory.find( { type: 'food' } );  
  
myCursor
```

```
var myCursor = db.inventory.find( { type: 'food' } );  
  
myCursor.forEach(printjson);
```

```
var myCursor = db.inventory.find( { type: 'food' } );  
  
var documentArray = myCursor.toArray();  
  
var myDocument = documentArray[3];
```

- Create an Index to Support Read Operations.
- Query Selectivity – the inequality operators `$nin` and `$ne` are not very selective
- Covering a Query – all the fields in the query and returned are part of the index
- Query Plans – query optimizer caches query plan

When Query Plan Revision happens?





As collections change over time, the query optimizer deletes the query plan and re-evaluates after any of the following events:

- » The collection receives 1,000 write operations.
- » The reIndex rebuilds the index.
- » You add or drop an index.
- » The mongod process restarts.



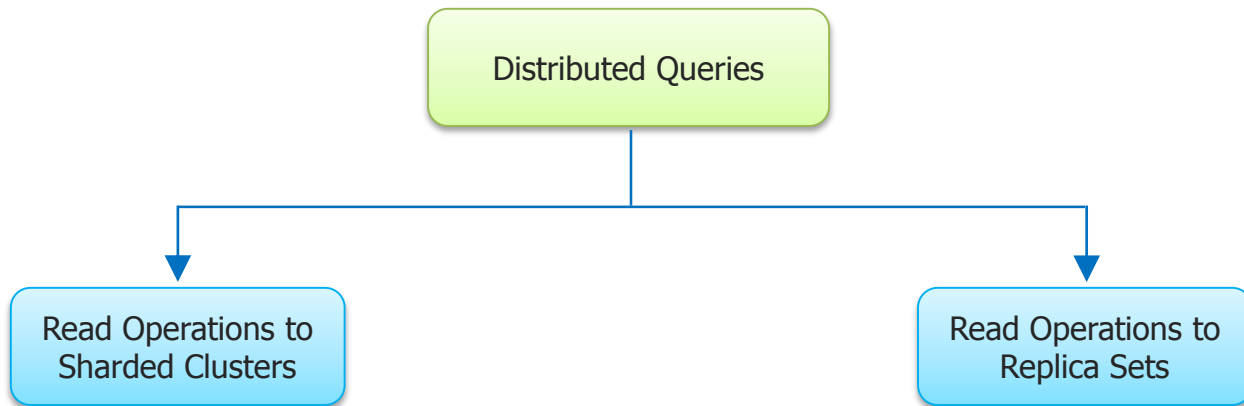
How do I isolate cursors from intervening write operations?



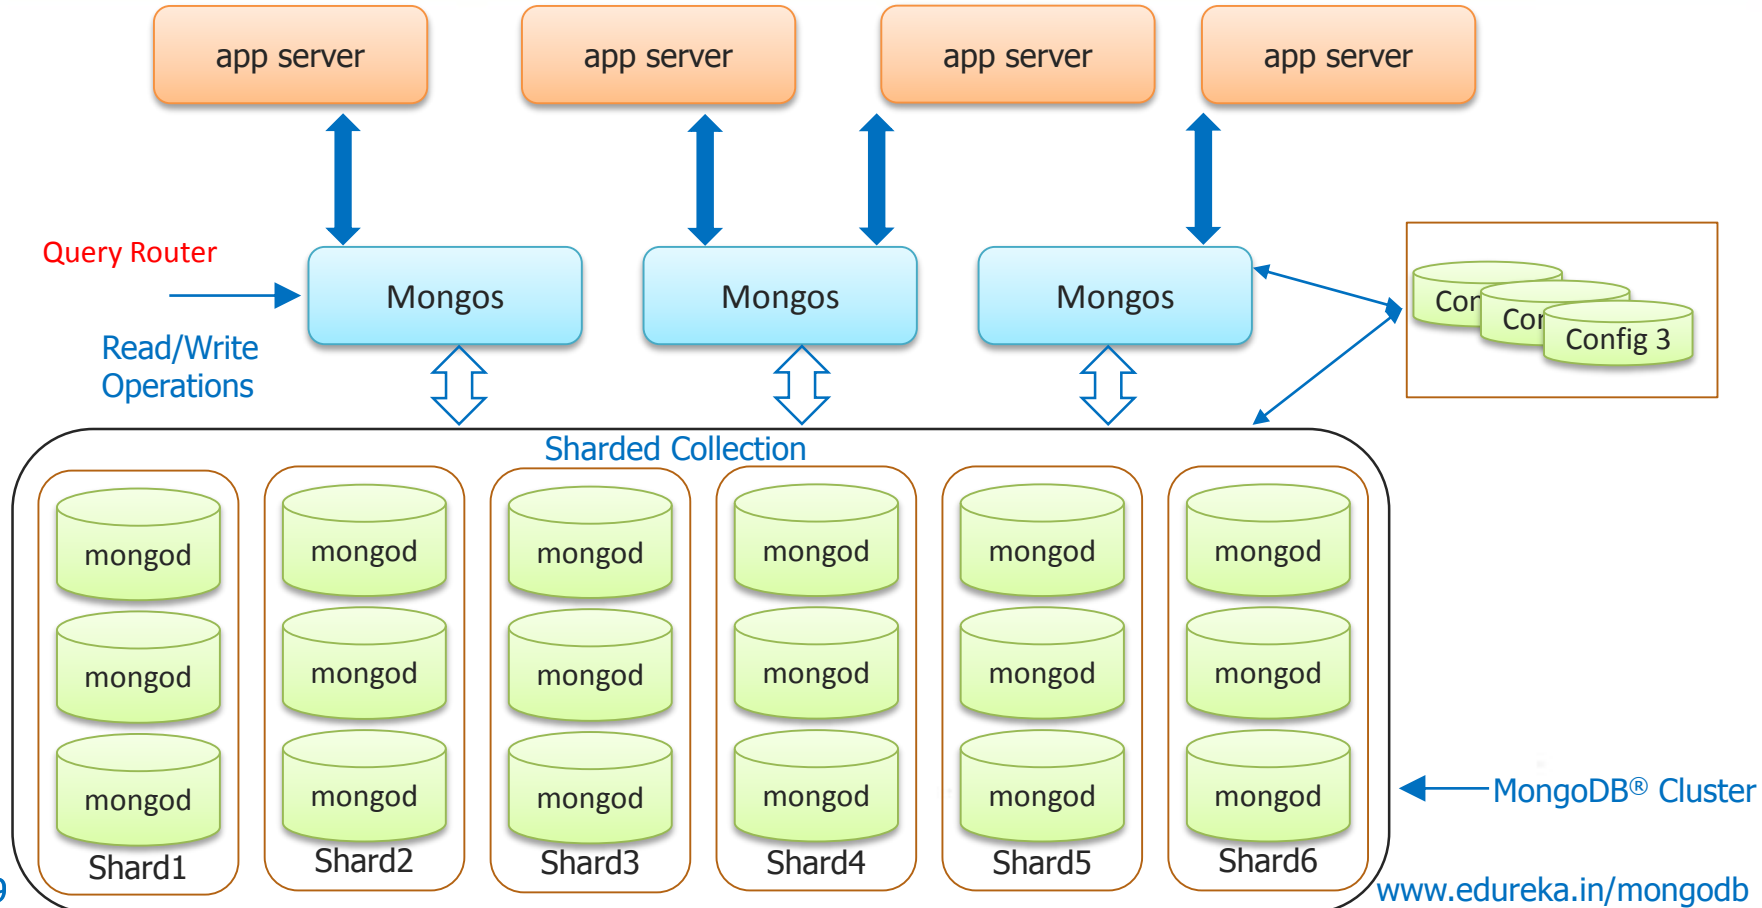
MongoDB® cursors can return the same document more than once in some situations.

You can use the `snapshot()` method on a cursor to isolate the operation for a very specific case.

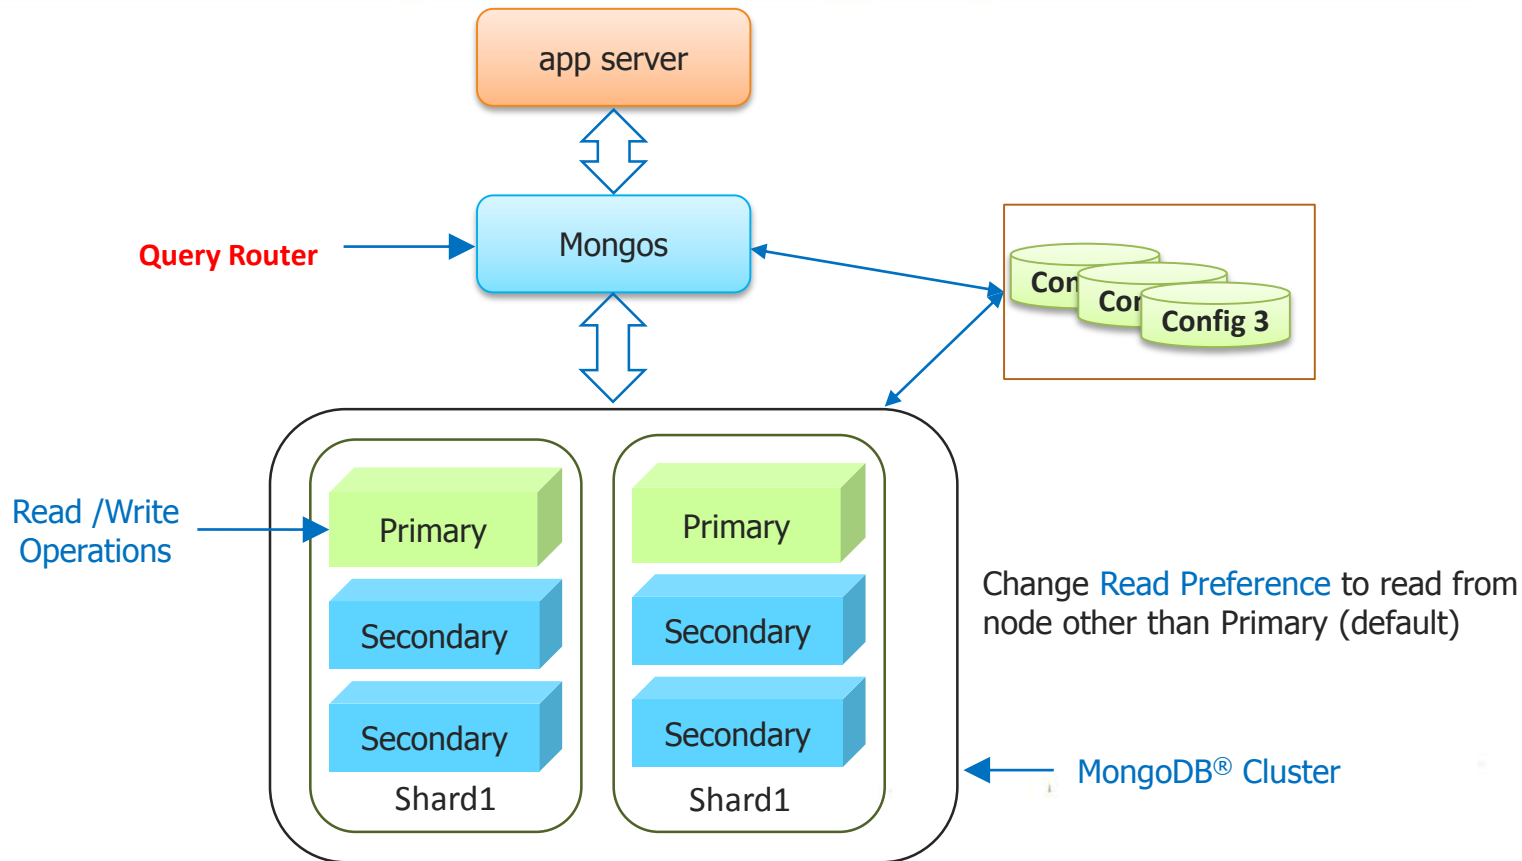




# Read Operations to Sharded Clusters



# Read Operations to Replica Set



# Write Operations Overview

Create in MongoDB® :

```
db . users . insert(  ← collection
{
    name : "sure", ← field value
    age: 26, ← field value
    status: "A" ← field value
}
)  ← Documents
```

Create in SQL:

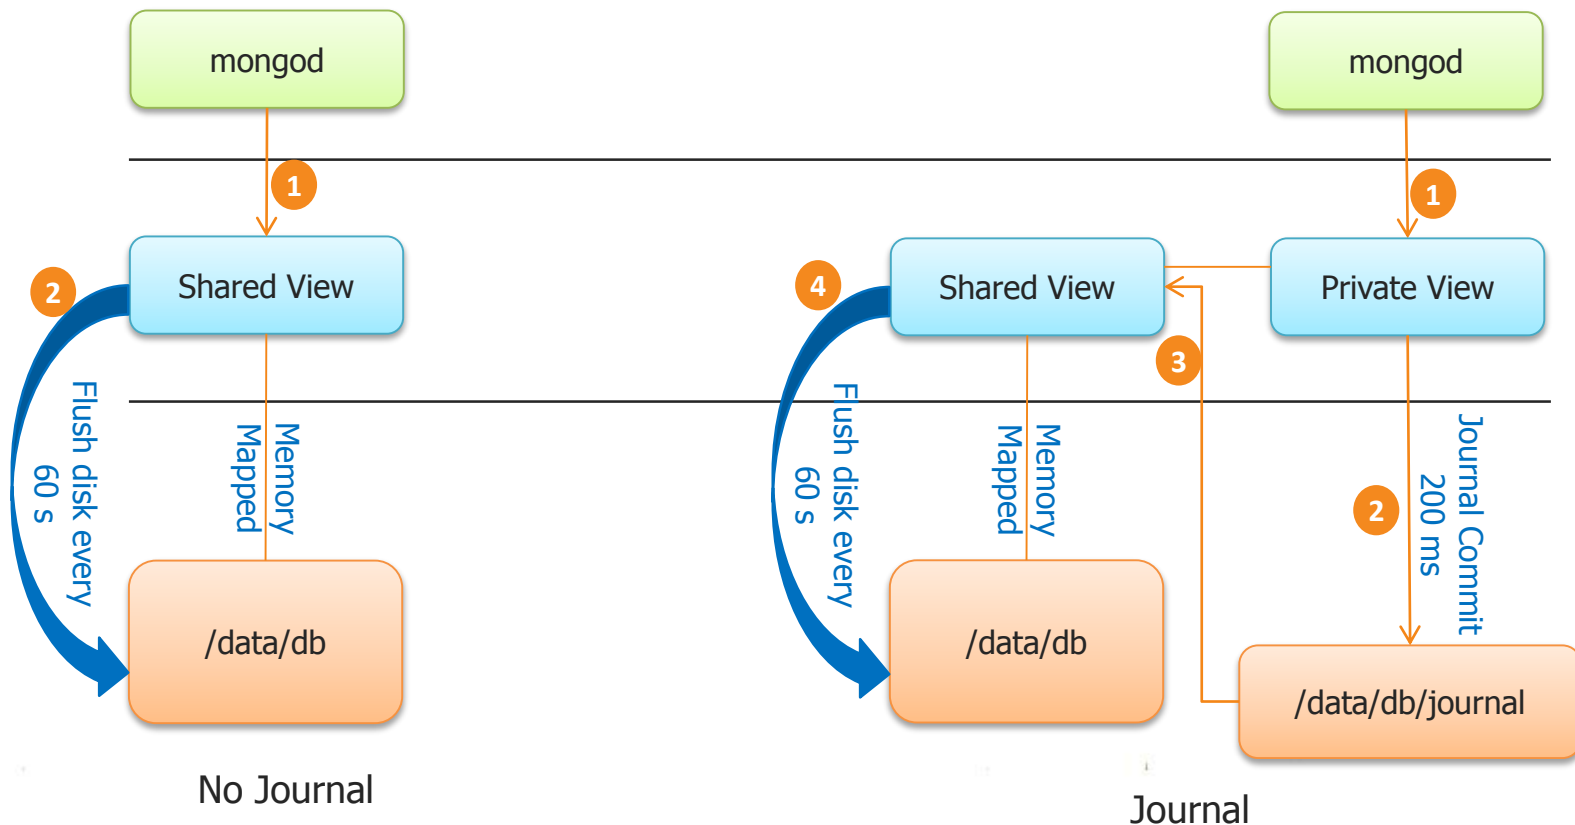
```
INSERT INTO users ← table
    (name, age, status) , ← columns
VALUES ("sure", "26", "A") , ← values/row
```

Update in MongoDB® :

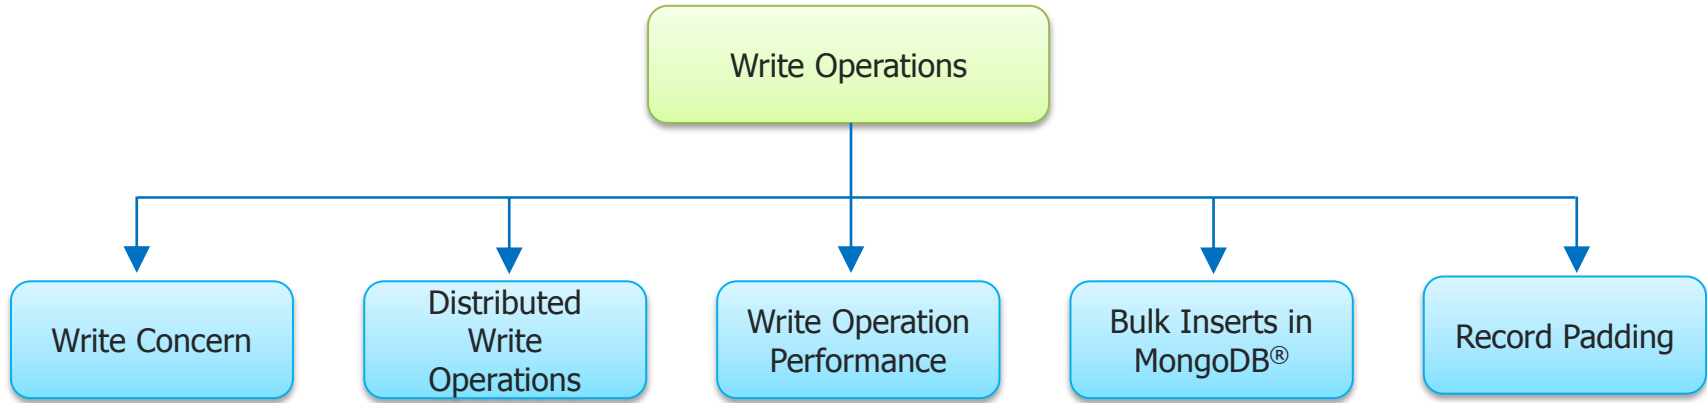
```
db . users . update(  ← collection
    { age : { $gt : 18 } } , ← update criteria
    { $set : { status: "A" } } , ← update action
    { multi : true } , ← update option
)
```

Update in SQL:

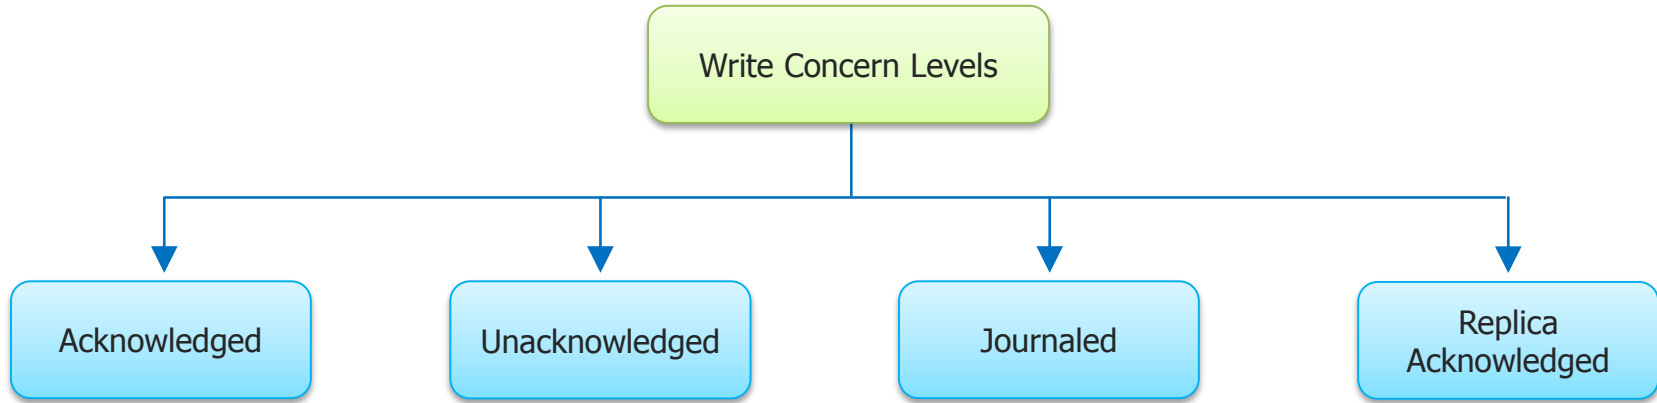
```
UPDATE users ← table
SET status = "A", ← update action
WHERE age > 18 ← update criteria
```



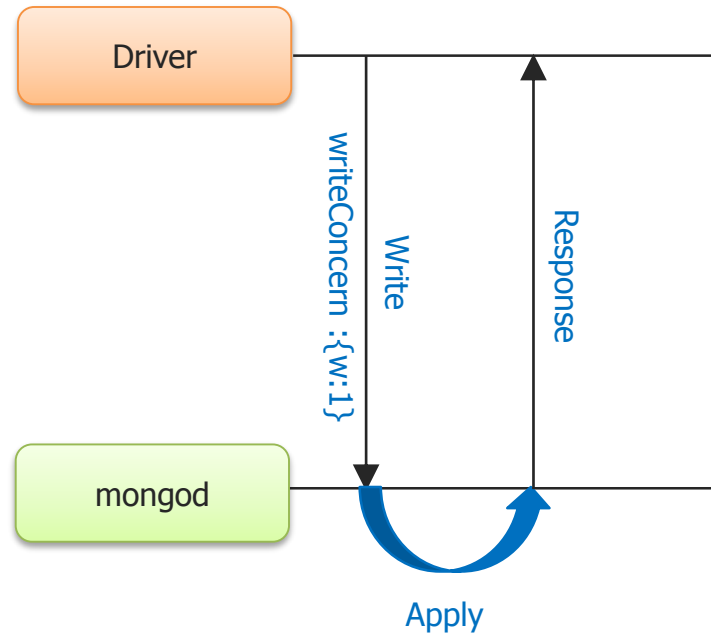




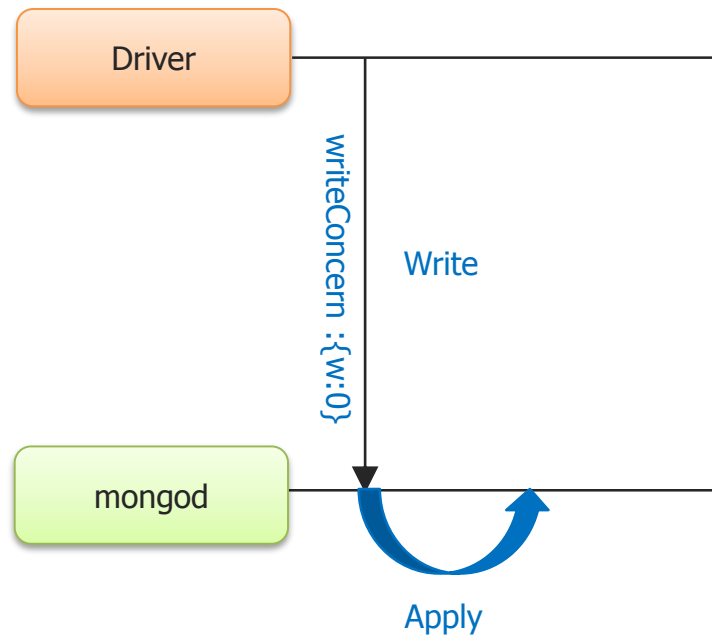
- Write concern describes the guarantee that MongoDB® provides when reporting on the success of a write operation.
- The strength of the write concerns determine the level of guarantee.
- When inserts, updates and deletes have a weak write concern, write operations return quickly.
- In some failure cases, write operations issued with weak write concerns may not persist.
- With stronger write concerns, clients wait after sending a write operation for MongoDB® to confirm the write operations.



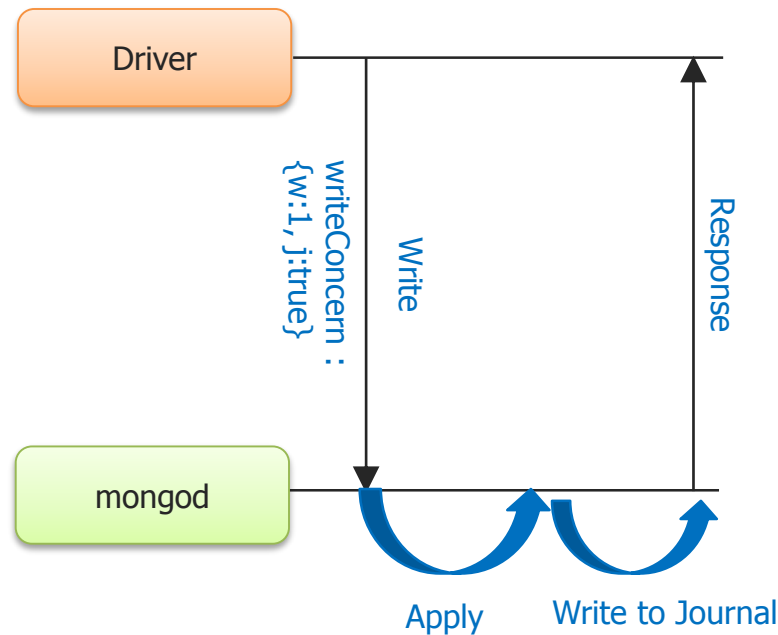
- With a receipt acknowledged write concern, the mongod confirms the receipt of the write operation.
- Acknowledged write concern allows clients to catch network, duplicate key, and other errors.
- To set acknowledged write concern, specify w values of 1 to your driver.
- MongoDB® uses acknowledged write concern by default, after the releases outlined in Default Write Concern Change.
- Write operation to a mongod instance with write concern of acknowledged the client waits for acknowledgment of success or exception.



- With an unacknowledged write concern, MongoDB® does not acknowledge the receipt of write operation.
- Unacknowledged is similar to errors ignored; however, drivers attempt to receive and handle network errors when possible.
- Write operation to a mongod instance with write concern of unacknowledged the client does not wait for any acknowledgment.

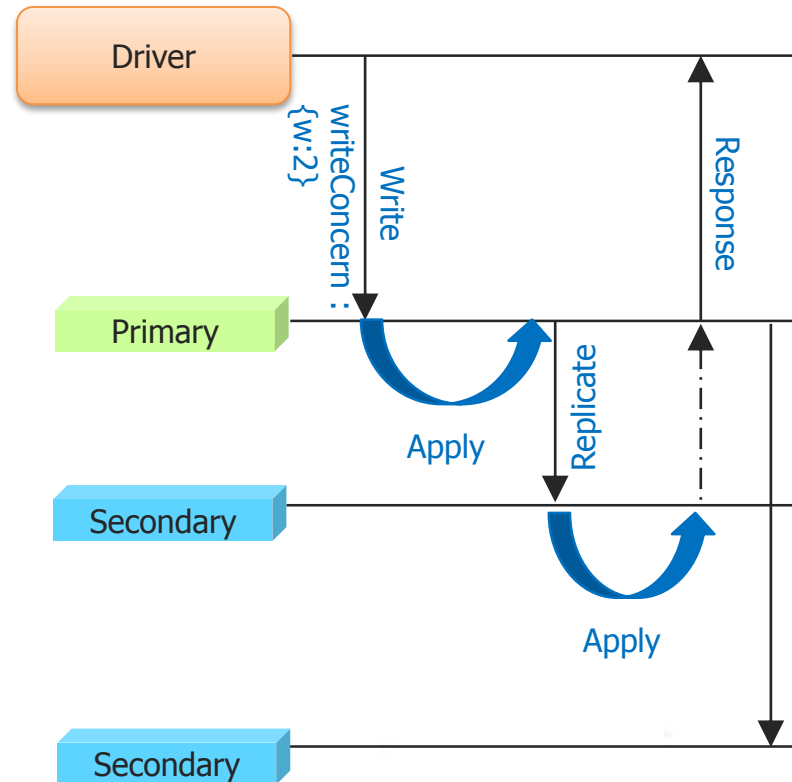


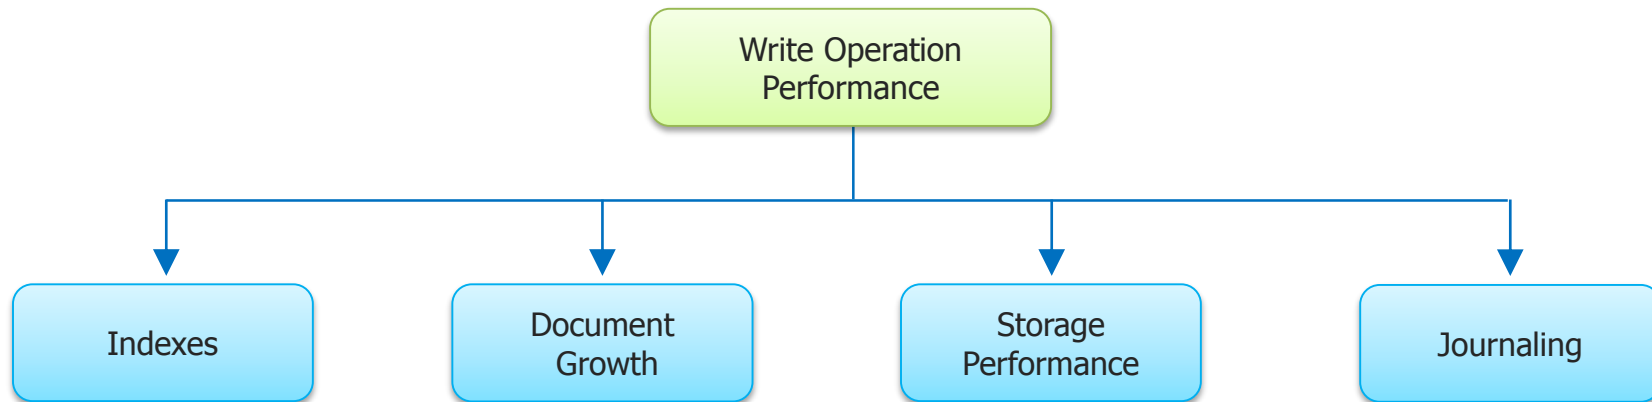
- With a journaled write concern, the mongod acknowledges the write operation only after committing the data to the journal.
- This write concern ensures that MongoDB® can recover the data following a shutdown or power interruption.
- To set a journaled write concern, specify w values of 1 and set the journal or j option to true for your driver.
- With a journaled write concern, write operations must wait for the next journal commit.
- To reduce latency for these operations, you can increase the frequency that MongoDB® commits operations to the journal.
- The mongod sends acknowledgment after it commits the write operation to the journal.



# Replica Acknowledged

- Replica sets add several considerations for write concern.
- Basic write concerns affect write operations on only 1 mongod instance.
- With replica acknowledged you can guarantee that the write operation propagates to the members of a replica set.
- To set replica acknowledged write concern, specify w values greater than 1 to your driver.







What is record padding?



Padding is a process to minimize document movements. Every document in MongoDB® is stored in a record which contains the document itself and extra space, or padding, which allows the document to grow as the result of updates.



<p><b>String:</b> Used to store the string data. String in MongoDB must be UTF-8 valid. (2)</p>	<p><b>Arrays:</b> This type is used to store arrays or list or multiple values into one key. (4)</p>	<p><b>Object:</b> This data type is used for embedded documents.</p>	<p><b>Code:</b> To store java script code into document. (13)</p>	<p><b>Time Stamp:</b> 64 bit value. 1<sup>st</sup> 32 bit seconds since epoch. 2<sup>nd</sup> 32 bit is incremental. (17) new Timestamp()</p>
<p><b>Integer:</b> For numerical value. Integer can be 32 bit or 64 bit depending upon your server. (16/18)</p>	<p><b>Boolean:</b> For a Boolean (true/ false) value. (8)</p>	<p><b>Double:</b> For floating point values. Default type. (1)</p>	<p><b>Regular expression:</b> To store regular expression. (11)</p>	<p><b>Null:</b> This type is used to store a Null value. (10)</p>
<p><b>Date:</b> For current date or time in UNIX time format. (9) Date() new Date() ISODate()</p>	<p><b>Min/ Max keys:</b> To compare a value against the lowest and highest BSON elements. (-1 / 127)</p>	<p><b>Symbol:</b> It's generally reserved for languages that use a specific symbol type. (14)</p>	<p><b>Object ID:</b> 12 byte BSON. 4 byte timestamp, 3 byte machine, 2 byte process id and last 3 byte counter. (7)</p>	<p><b>Binary data:</b> To store binary data. (5)</p>

- Insert Documents
- Query Documents
- Limit Fields to Return from a Query
- Iterate a Cursor in the mongo Shell
- Analyze Query Performance
- Modify Documents
- Remove Documents

LIVE DEMO

What is two phase commit?








This is used to ensure immediate consistency and reliability.

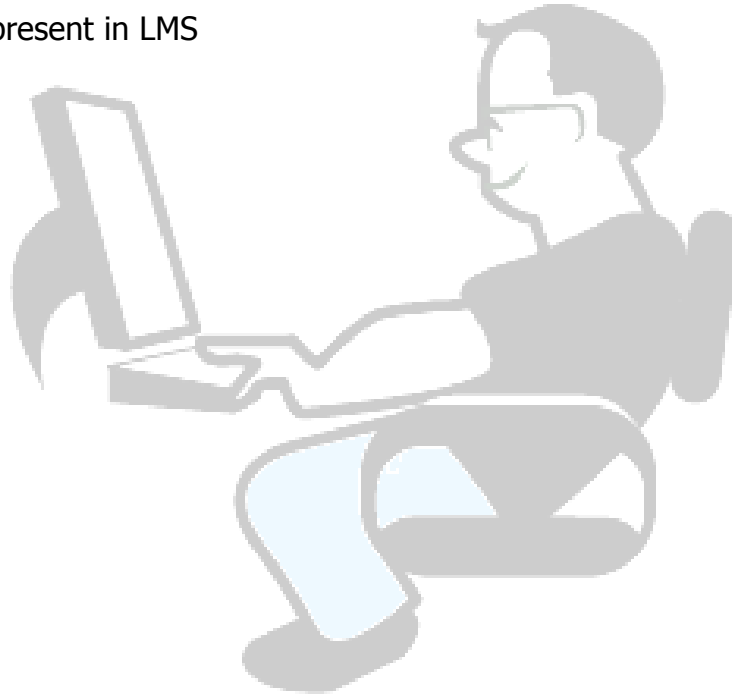


- How to iterate a Cursor in the mongo shell?
- How to close Inactive Cursors?
- What is Tailable Cursor?
- How can you create an Auto-incrementing Sequence Field?
- What is Record Padding?
- How many methods are there to insert document in Collection?
- What is \_id field in MongoDB® ?
- What all the operators are there in MongoDB®
- What is capped collection?
- What is CRUD?





-  Generate Test Data on MongoDB® Database
-  Execute all Module2 Script present in LMS
-  Read Module 2 FAQ
-  Attempt Module 2 Quiz
-  Complete assignment



# Agenda for Next Class

- Data Modeling Concepts
- Type of Data Modeling
- Why Data Modeling ?
- Data Modeling Approach
- Analogy Between RDBMS & MongoDB® Data Model
- MongoDB® Data Model
- Challenges for Data Modeling in MongoDB®
- Data Model Examples and Patterns
- Data Model References
- Use Case of Data Modeling



Your feedback is important to us, be it a compliment, a suggestion or a complaint. It helps us to make the course better!

Please spare few minutes to take the survey after the webinar.

Thank you!

